

# SRECHARAN SELVAM

☎ +1 917-495-3410 ✉ [sselvam@andrew.cmu.edu](mailto:sselvam@andrew.cmu.edu) 🔗 [linkedin.com/srecharan](https://www.linkedin.com/srecharan) 🌐 [Srecharan.com](https://Srecharan.com) 🐙 [github.com/Srecharan](https://github.com/Srecharan)

## Education

### Carnegie Mellon University

*Master of Science, Research: Machine Learning*

Relevant Coursework: Machine Learning, Deep Learning, Advanced Computer Vision, Visual Recognition

May 2025

Pittsburgh, PA

GPA - 4/4

## Experience

### Kantor Lab, Carnegie Mellon University [GitHub ↗]

Aug 2023 – Present

*Graduate Research Assistant*

- Developed self-supervised learning pipeline combining RAFT-Stereo depth estimation, YOLOv8 segmentation & geometric algorithms, eliminating **100%** manual annotation requirements to train GraspPointCNN for robotic leaf manipulation.
- Trained attention-based GraspPointCNN using MLflow to track **60+** model experiments for grasp point optimization.
- Parallelized 1.5M-pixel depth-to-3D projection via CUDA kernels for efficient real-time 3D reconstruction (**150→30ms**).
- Compiled vision models into TensorRT engines with FP16 precision to boost inference throughput from **20 to 27 FPS**.
- Deployed Docker-containerized grasping system to a 6-DOF robot, achieving **82%** leaf grasp success rate in field tests.

### Hanon Systems [GitHub ↗]

Jan 2023 – Jun 2023

*Machine Learning Engineer Intern*

- Built a GRU-based model trained on custom dataset with EKF for real-time 3D hand gesture (dynamic/static) recognition in AR system (Unity Engine), allowing **100+** automotive technicians to safely simulate HVAC assembly procedures.
- Optimized inference using CUDA kernels and ONNX quantization, cutting latency by **33%** & memory footprint by **50%**.
- Exposed gesture and depth modules via Flask REST APIs for modular model serving to downstream applications.
- Containerized complete pipeline by utilizing Docker for consistent, scalable deployment across **3** training centers.

### Vee Ess Engineering [GitHub ↗]

Jul 2022 – Dec 2022

*Computer Vision Engineer Intern*

- Engineered distributed Apache Spark pipeline to boost recyclable material recovery on high-speed conveyors by handling terabytes of multi-camera footage and auto-annotating frames with Mask R-CNN to create **43,000+** segmented images.
- Streamlined dataset storage & versioning via AWS S3 to reduce I/O overhead by **24%** and accelerate training iterations.
- Integrated custom-trained YOLOv5 model across real-time camera streams yielding **96% mAP@[0.5:0.95]** with **<15ms** latency, enabling conveyor speed modulation based on detection density to cut manual sorting labor by **6+ hrs/week**.

## Projects

### VLM-Based Tool Recognition System for Industrial Safety Applications [GitHub ↗]

Jan 2025 – May 2025

- Fine-tuned Qwen-2.5-VL-7B & LLaMA-3.2-11B-V using LoRA on custom dataset via LLM-guided prompt engineering (**8K** images, **29K** annotations) for real-time multi-modal industrial tool recognition and safety guidance generation.
- Built a RAG pipeline using LangChain and Pinecone to ground tool-specific information, reducing hallucinations by **55%**.
- Implemented RLHF (GRPO) on AWS SageMaker to optimize preference learning on paired responses for VLM alignment.
- Orchestrated LLM-based evaluation pipeline (OpenAI API) with Kubernetes, scoring **4K+** outputs for **8** model variants.

### Multi-Model Stock Prediction with NLP and Automated Trading [GitHub ↗]

Oct 2024 – Feb 2025

- Spearheaded distributed training infrastructure using data parallelism across 4x V100 GPUs to train ensemble ML models (bidirectional LSTM + XGBoost with **35+** features) for algorithmic trading system across multiple timeframes.
- Created Kafka streaming pipeline to ingest **9K** financial events/day from multiple APIs for low-latency trading decisions.
- Automated Apache Airflow workflows managing FinBERT sentiment analysis, boosting prediction accuracy by **~5%**.
- Designed automated trading system with CI/CD model retraining and Tradier API execution, delivering **58.5%** win rate.

### Bird Image Generation using GANs, VAEs, and Diffusion Models [GitHub ↗]

Apr 2024 – Jul 2024

- Trained WGAN-GP with gradient penalty,  $\beta$ -VAEs, and diffusion models on CUB-200-2011 dataset for image synthesis.
- Developed custom loss functions to optimize image quality, achieving best performance with WGAN-GP (**33.07 FID**).
- Accelerated diffusion inference **10x** (1000 to 100 time-steps) via DDIM sampling while maintaining same image quality.

## Skills

**Languages & Frameworks:** Python, C++, SQL, PyTorch, TensorFlow, OpenCV, scikit-learn, Transformers, ONNX, Git  
**ML Training:** RAG, RLHF (GRPO, DPO), SFT, PEFT, LoRA, VLM/LLM Fine-tuning, Distributed Training, LangChain  
**Infrastructure:** AWS (SageMaker, EC2, S3), GCP, Docker, Kubernetes, TensorRT, CUDA, MLflow, Kafka, Spark, Flask

## Publications

Srecharan Selvam, Abhisesh Silwal, George Kantor "Self-Supervised Learning for Robotic Leaf Manipulation: A Hybrid Geometric-Neural Approach", [arXiv:2505.03702](https://arxiv.org/abs/2505.03702), Under review at ICCV 2025. [Project Page ↗] [PDF ↗]